



# The Beauty and Joy of Computing

## Lecture #07 Concurrency

UC Berkeley  
Teaching Assistant  
Aaron Baker



### **“KOOMEY’S LAW” – EFFICIENCY 2X EVERY 18**

Prof Jonathan Koomey looked at 6 decades of data and found that energy efficiency of computers doubles roughly every 18 months. This is even more relevant as battery-powered devices become more popular. Restated, it says that for a fixed computing load, the amount of battery you need drops by half every 18 months. This was true before transistors!





# Concurrency: Definition & Examples

---

Definition: A property of computer systems in which several computations are executing simultaneously, and potentially interacting with each other.

## Examples:

- Mouse cursor movement while Snap! calculates.
- Screen clock advances while typing in a text.
- Busy cursor spins while browser connects to server, waiting for response



# Concurrency & Parallelism

## Intra-computer

- Today's lecture
- Multiple computing “helpers” are cores within one machine
- Referred to as “multi-core”




## Inter-computer

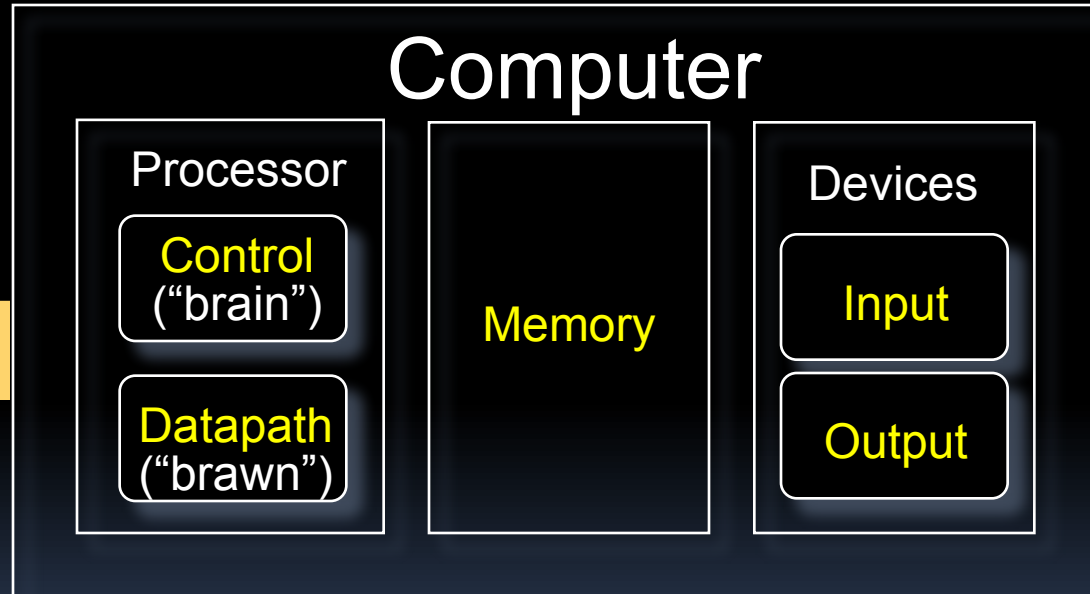
- Multiple computing “helpers” are different machines
- Referred to as “distributed computing”



# Anatomy: 5 components of any Computer




John von Neumann  
invented this  
architecture

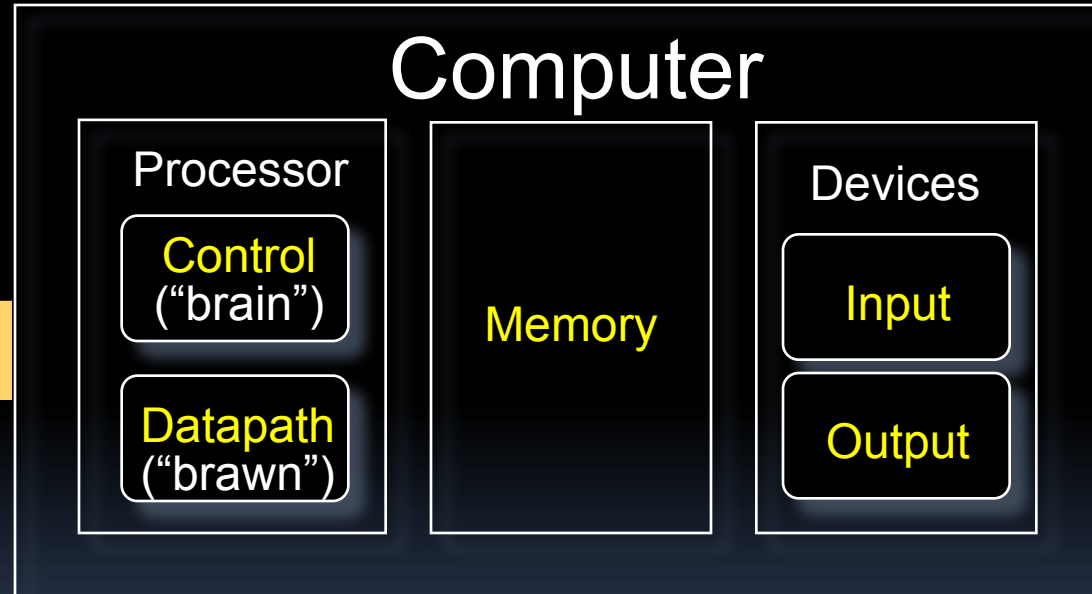


# Anatomy: 5 components of any Computer

[www.yellkey.com/fly](http://www.yellkey.com/fly)



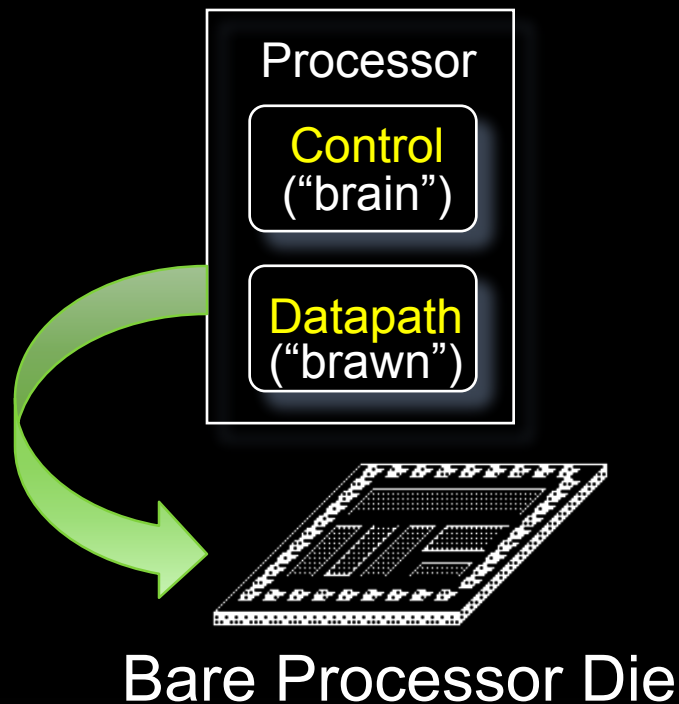
- a) Control
- b) Datapath
- c) Memory
- d) Input
- e) Output




What causes the most headaches for SW and HW designers with multi-core computing?

Garcia

# But what is INSIDE a Processor?



- Primarily Crystalline Silicon
- 1 mm – 25 mm on a side
- 2009 “feature size” (aka process)  
~ 45 nm =  $45 \times 10^{-9}$  m  
(then 32, 22, and 16 [by yr 2013])
- 100 - 1000M transistors
- 3 - 10 conductive layers
- “CMOS” (complementary metal oxide semiconductor) - most common





Chip in Package

- Package provides:
  - spreading of chip-level signal paths to board-level
  - heat dissipation.
- Ceramic or plastic with gold wires.

# Moore's Law

Predicts: 2X Transistors / chip every 2 years

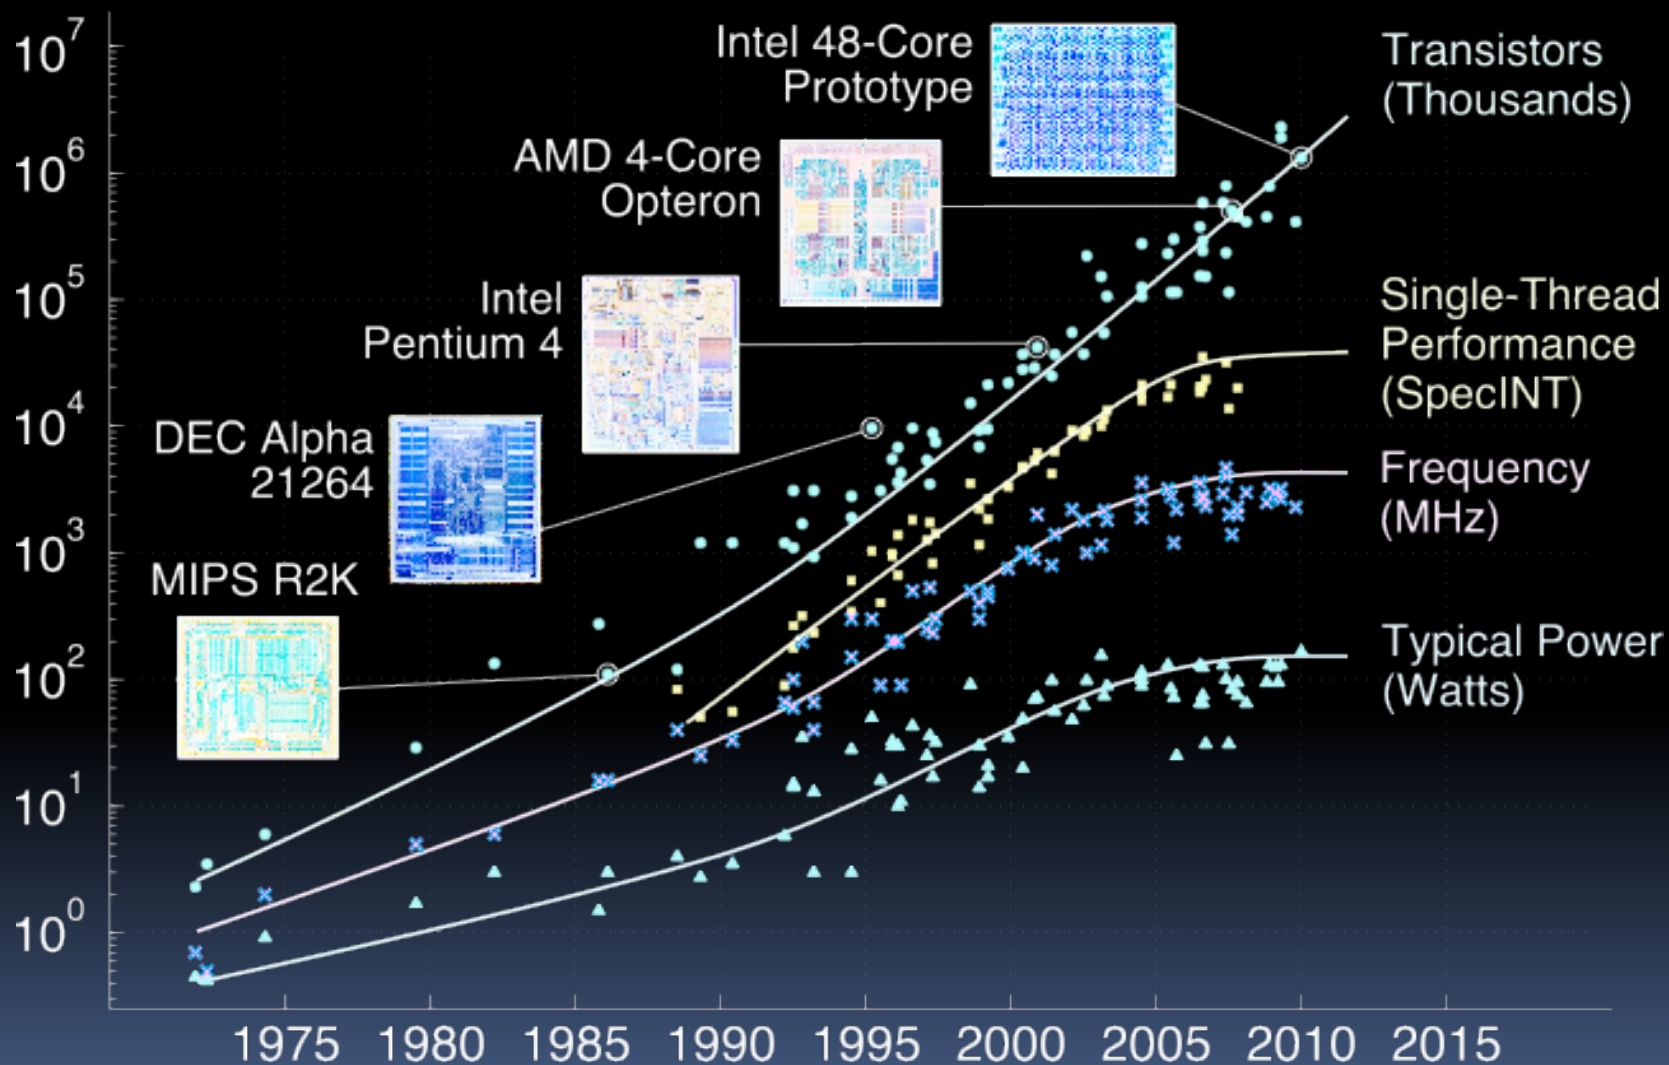


Gordon Moore  
Intel Cofounder  
B.S. Cal 1950!

[www.yellkey.com/security](http://www.yellkey.com/security)



# Moore's Law and related curves




Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond

Garcia



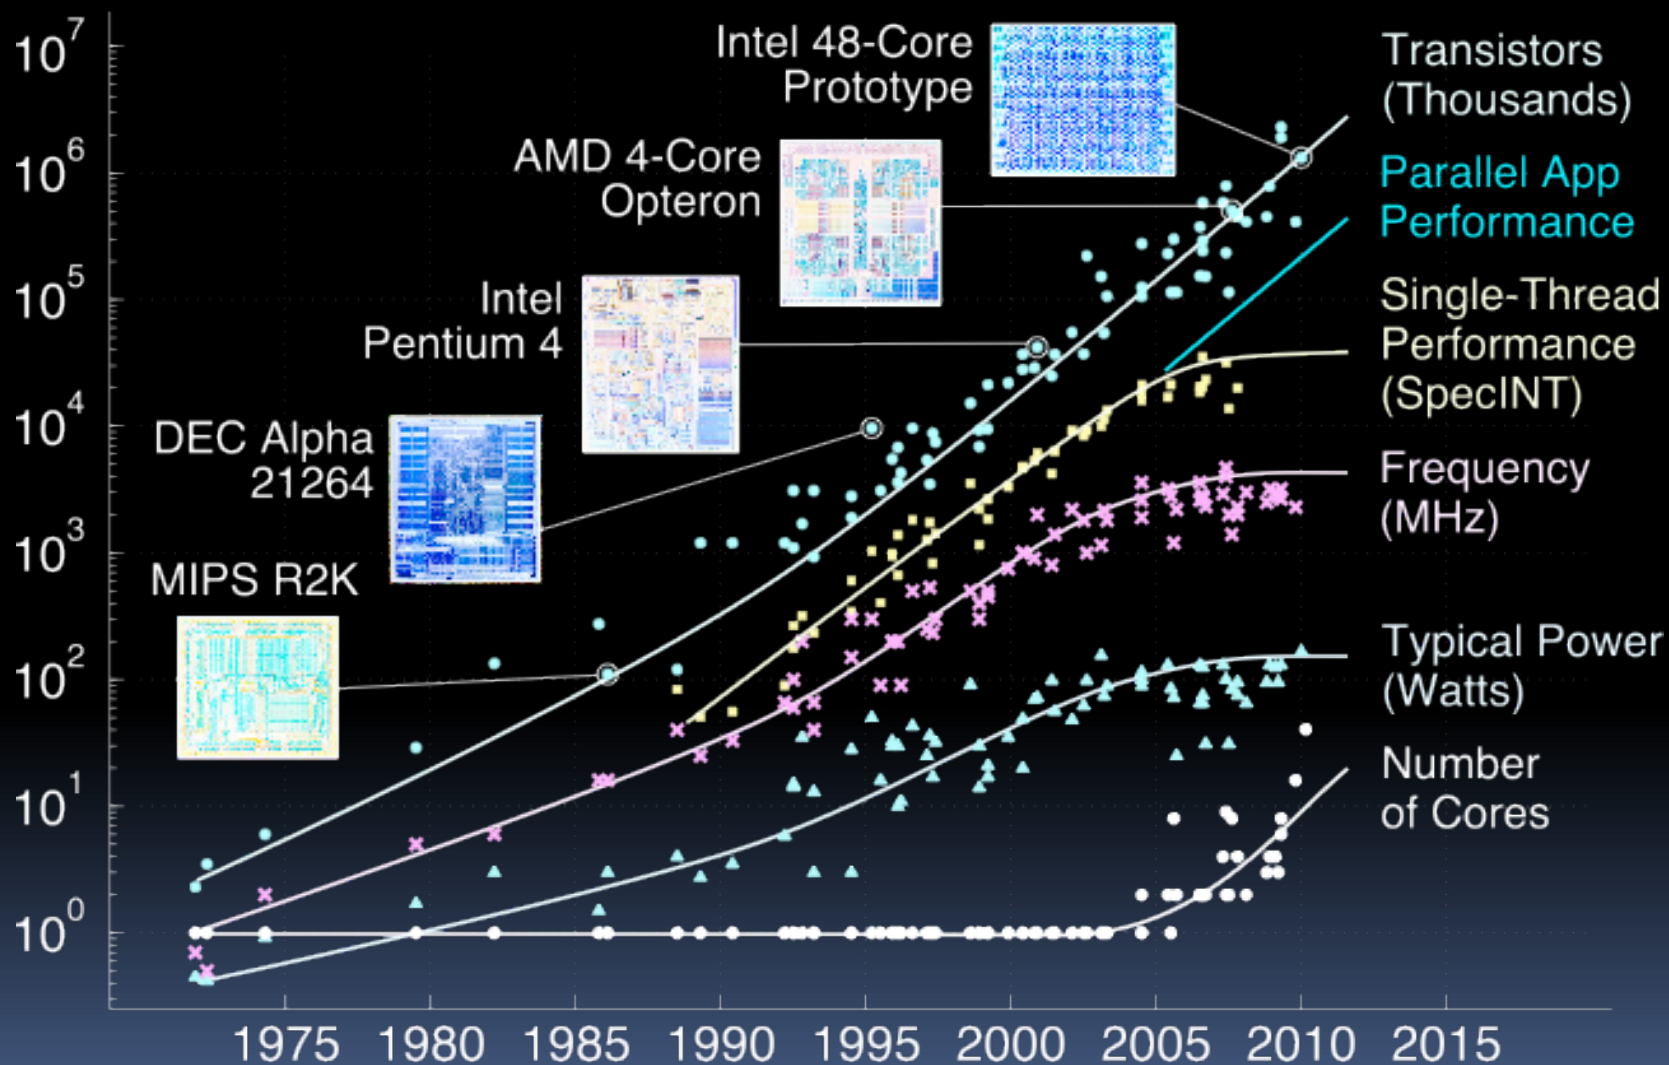
# Power Density Prediction circa 2000



Source: S. Borkar  
(Intel)

Garcia

# Moore's Law and related curves




Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond

Garcia

# Parallel Computing: Threads

- A **Thread** stands for “thread of execution”; it is a single stream of instructions
  - A program / process can **split** or **fork** itself into separate threads, which can (in theory) execute simultaneously.
  - An easy way to describe/think about parallelism
- With a single core, a single CPU can execute many threads by **Time Sharing**



■ Thread<sub>0</sub>  
■ Thread<sub>1</sub>  
■ Thread<sub>2</sub>

- **Multithreading** is running multiple threads through the same hardware



# Speed: Making a Pizza


---

- Tasks:
  - Make the dough: 10 minutes
  - Make the sauce: 10 minutes
  - Prepare the toppings: 10 minutes
  - Assemble the pizza: 10 minutes
  - Bake the pizza: 40 minutes
- How long would this take with just one person?
- With two people? Three people? Four?
- Will adding more people always speed up the process?



# Speedup Issues : Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains



- $s$  is serial fraction of program,  $P$  is # of cores (was processors)

- Amdahl's law:

$$\text{Speedup}(P) = \text{Time}(1) / \text{Time}(P)$$

$$\leq 1 / (s + [(1-s) / P]), \text{ and as } P \rightarrow \infty$$

$$\leq 1 / s$$

- Even if the parallel portion of your application speeds up perfectly, your performance may be limited by the sequential portion

# Speedup Issues : Overhead

- Even assuming no sequential portion, there's...
  - Time to think how to divide the problem up
  - Time to hand out small “work units” to workers
  - All workers may not work equally fast
  - Some workers may fail
  - There may be contention for shared resources
  - Workers could overwriting each others' answers
  - You may have to wait until the last worker returns to proceed (the slowest / weakest link problem)
  - There's time to put the data back together in a way that looks as if it were done by one



# Life in a multi-core world...

- This “sea change” to multi-core parallelism means that the computing community has to rethink:

- a) Languages
- b) Architectures
- c) Algorithms
- d) Data Structures
- e) All of the above



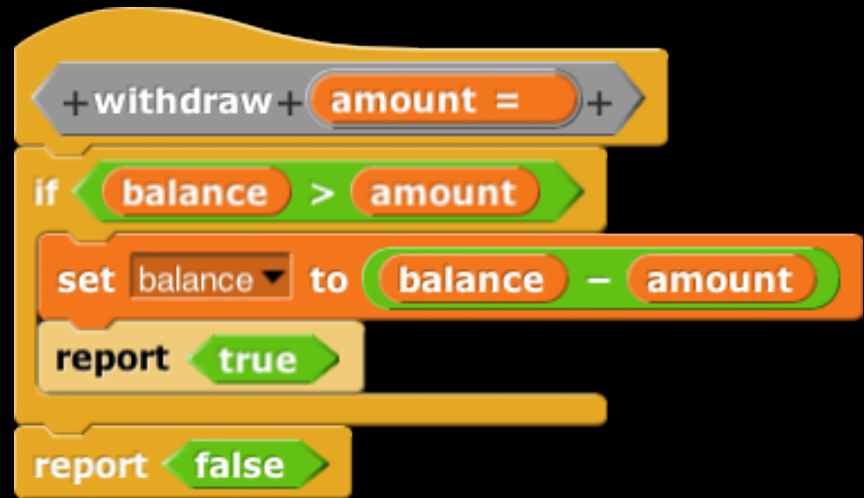
[www.yellkey.com/still](http://www.yellkey.com/still)



# Concurrency Issues

- What if two people were calling withdraw at the same time?

- ☐ Assume balance=\$100 and two people withdraw \$75 each (same account)
- ☐ Computer will try to parallelize this process by executing each call to “withdraw” on a separate core



- Can anyone see what the problem could be?
  - Problem: There is only one copy of “balance” in memory. If both threads access “balance” simultaneously, people will withdraw more money than they actually have!
  - Bigger problem: Program may or may not work properly





# Race Condition


---

- Two (or more) scripts are running at the same time, BUT **we don't know what order they will be run in!** This is called a **race condition**.
- Very common in many programming languages. **This will not happen in Snap!** because it will swap between running script A and script B and won't let both of them run at once. (Phew!)
- ...but if you have some explicit calls to “wait random secs” in your code & use these for timers, it can happen!



# Possible Outputs?

- We want this code to draw a cute winky-face, but there's a problem with parallelizing it!
- Which of the following outputs is not possible?
  - A) No output
  - B) Just left eye
  - C) Just left eye and smile
  - D) Just right eye and smile
  - E) Whole face



# Another concurrency issue... **deadlock!**

---

- **Deadlock**: Threads are waiting on each other before acting, no action
  - Ex: You see an acquaintance walking on campus, don't want to say hi first. Your acquaintance doesn't want to say hi first either. You wait on your acquaintance to say hi first, your acquaintance waits on you, no one says hi
- **Livelock** also possible: Threads keep taking action without waiting for each other, no progress
  - Ex: Passing someone in a narrow hall, keep moving out of each other's way and blocking each other



# Summary

- “Sea change” of computing because of inability to cool CPUs means we’re now in multi-core world
- This brave new world offers lots of potential for innovation by computing professionals, but challenges persist

